

# Robust Timeseries Analysis in the Context of Multiple Asynchronous Input Channels

Ryan Michael  
kerinin@gmail.com

March 3, 2009

## Abstract

Spectral Analysis is a common form of analysing timeseries data. We propose a method of spectral analysis based on sets probability distributions generated from subsets of the observed data. By decomposing the observations into a set of observed behaviors, the spectral analysis can be used more accurately to predict the timeseries' behavior.

Decomposed into the ??? domain?

## 1 Introduction

### 1.1 Existing Work

### 1.2 General Overview

The goal is to create a method of statistical inference capable of processing timeseries data which is both multi-variate and exhibits different behaviors at different times. This type of data is common, and developing a robust method of analysis has applications in many domains. The general approach is to create a series of estimates using subsets of the observed data, and to then combine these estimates in an intelligent manner which captures the relevance of each estimate to the current prediction task. By using a set of 'typical' estimates, we are able to reduce the computational demands of the system, as each estimate is a condensed representation of the data from which it was derived. This approach also allows us to reduce data redundancy by only using distinct estimates.

The most basic operation used in this system is the estimation of probability densities. Based on a set of observations drawn from some random process, we generate an estimate of the underlying probability distribution. This estimate tells us the probability of each point in the input space being observed. Areas

of the input space in which a dense set of observations are observed are given high probability, while areas of the input space with few observations are given low probability. This basic operation, in combination with some basic laws of statistics allow us to build the full inference system.

We assume that observations are pulled from multiple independent sources, all of which respond to some underlying phenomena. For instance one set of observations could be from a microphone and another from a light detector. We do not know how the inputs are related or to what extent their behavior changes over time. For example one input could be a steady sine wave and another could be based on the decay of a radioactive isotope. In the former case previous observations of the source are useful, in the latter they're not. Alternately two inputs could be light detectors in the same room or they could be on different continents; in the former their input would be highly similar, in the latter not as much.

Our general strategy has three phases; generating estimates, correlating estimates, and applying estimates to a given prediction task.

### 1.3 Generating Estimates

Estimates are generated by picking a time window at random and only dealing with observations which take place in that window. Using these observations we create an estimate of the probability distribution underlying the observations (this distribution would include time as a variable). This process is repeated until we feel we have a reasonable sample of the system as a whole, at which time we can start to correlate the estimates.

### 1.4 Correlating Estimates

The goal of correlating estimates is to determine the relationships between estimates at different times and from different sources. Estimates are correlated by treating them as variables whose value for a given time window is determined by the extent to which the observed data corresponds to the estimate. For each time window there exist a set of estimates which have some value describing their accuracy at predicting the observed value. We can treat the accuracy measurement of each estimate as a multi-dimensional point, each dimension determined by an estimate's accuracy. Using a set of these points taken at different times, we create a probability distribution estimate. The domain of this probability distribution has the same dimensionality as the number of estimates we have generated. This probability density allows us to predict the probability of an estimate in of one source based on the probability of an estimate in of another source because frequent combinations of accuracy values will have higher probability than other combinations. This 'correlation density' also tells us which estimates are most commonly observed - information we can use in conjunction with estimate similarity to determine which estimates to use and which to discard.

## 1.5 Making Predictions

Once the estimates have been correlated, we are able to generate predictions. For simplicity, we'll assume that the first two phases occur on a set of 'training' data which is representative of the underlying data, while prediction takes place continuously using new sets of observations which occur in some time window. We make predictions for a given source by combining the existing estimates we have for that source based on their accuracy at predicting the given observations. Because we have determined the correlation between estimate accuracy of different sources, we can use other sources to refine our confidence in each estimate of the given source; the influence of estimates of the given source which do not correspond to a likely 'point' in the correlation density are suppressed while the influence of estimates which correspond to likely points in the correlation density are enhanced.

## 2 Formal Description

### 2.1 Problem Setting

We begin with a hidden random variable

$$X = (\Omega, \mathcal{F}, \mathcal{P})$$

Our knowledge of  $X$  comes from a set of independent sources which we treat as random variables generated by  $X$ :

$$\begin{aligned} X^n : \Omega &\mapsto \Omega^n \in \mathbb{R}^d \times \mathbb{R}_+ \\ X^n &= (\Omega^n, \mathcal{F}^n, \mathcal{P}^n) \\ \mathbf{X} &= [X^0, \dots, X^c] \end{aligned}$$

We refer to each of these sources as a channel, and refer to each channel as the  $i^{\text{th}}$  element of the set  $\mathbf{X}$ :

$$C^n = [\Omega^n, \mathcal{F}^n, \mathcal{P}^n]$$

For each channel we are given a set of  $\ell$  observations of dimension  $d$ , each with a time value:

$$X^n = [(x_0^n, t_0^n), \dots, (x_\ell^n, t_\ell^n)] \in \mathbb{R}^d \times \mathbb{R}_+$$

We pick at random a set of  $z$  time windows, each defined by a starting time  $t$  and a duration  $\theta$ :

$$\mathcal{T} = [(t_0, \theta_0), \dots, (t_z, \theta_z)]$$

We assume that the probability of  $X$  is a time-dependent mixture of some set of distributions  $[f_0, \dots, f_w]$  whose influence is unknown and changes over time. We can refer to the weighted mixture which corresponds to a given time window  $(t, \theta)$  as:

$$\mathcal{P}(t, \theta) = \sum_i \delta_i \cdot f_i \quad (1)$$

where  $\delta_n$  is the weight corresponding to  $f_n$ . Finally, we assume that a similar mixture of densities can be determined for each channel:

$$\mathcal{P}^n(t, \theta) = \sum_i \delta_i^n \cdot f_i^n \quad (2)$$

## 2.2 Single Channel Estimation

We begin by considering the case where only one channel exists, so for now we will omit the superscript and refer to  $X^n = (\Omega^n, \mathcal{F}^n, \mathcal{P}^n)$  as  $X = (\Omega, \mathcal{F}, \mathcal{P})$ . Our task is to determine both a set of underlying distributions  $[f_0, \dots, f_w]$  and their relative influence over time. We begin by restricting our attention to the set of observations  $S_{t,\theta}$  which fall into the window  $(t, \theta)$ . We estimate a set of probability distributions  $\Phi$  using some function  $\psi_E$  based on the sets of observations in different windows:

$$S_{t,\theta} = \left[ \left( x_i, \frac{t_i - t}{\theta} \right) \mid x_i \in X, t \leq t_i < t + \theta \right] \quad (3)$$

$$\psi_E(S_{t,\theta} : \alpha_E) \mapsto \phi_n \simeq \mathcal{P}(t, \theta) \quad (4)$$

$$\Phi = [\phi_0, \dots, \phi_p]$$

The estimation algorithm  $\psi_E$  produces an estimate  $\phi_n$  of the probability distribution  $\mathcal{P}(t, \theta)$  for the random variable  $X$ . Multiple algorithms exist to accomplish such density estimations this, so details will be omitted.

## 2.3 Single Channel Correlation

Given a set of estimations, we need a method to relate each estimate's relevance to observations in different time windows. We use the Shannon entropy of the set of observations in a given time window with respect to a given estimation as a measure of the estimate's relevance to the observed data. A set of data points which conform to the predicted values of an estimate will have lower entropy than a set of data points which diverge from the same estimate.

$$H_n(S_{t,\theta}) \mapsto \delta_{\phi_n,t,\theta} = \sum_{(x,t) \in S_{t,\theta}} -\phi_n(x,t) \log \phi_n(x,t) \quad (5)$$

$$\Delta_n = [\delta_{\phi_n,t,\theta} \mid (t,\theta) \in \mathcal{T}]$$

The entropy for each estimate over different time windows defines a new random variable which we will refer to as  $H_n$ . We will compute a density estimate using  $\psi_C$  for  $H_n$  as we have with the input data previously:

$$H_n = (\Omega^H, \mathcal{F}^H, \mathcal{P}_n^H), \quad \Omega^H \in \mathbb{R}, \mathcal{F}^H = \mathcal{B}(\mathbb{R}) \quad (6)$$

$$\psi_C(\Delta_n, \alpha_C) \mapsto \varphi_n \quad (7)$$

$$\varphi_n(\delta) \simeq P(H_n = \delta) \quad (8)$$

## 2.4 Single Channel Prediction

The derivation algorithm produces predictions of  $X$  over a time window  $(t, \theta)$  based on a set of observations which occur in that time window  $S_{t,\theta}$ . Rather than using the estimation algorithm  $\psi_E$  to predict these values, however, the derivation algorithm  $\psi_D$  predicts the probability distribution using bayesian modification of existing estimates. To accomplish this, we use the conditional probability of the entropy value.

$$P(X = (x, t) \mid H_n(S_{t,\theta}) = \delta) = \frac{P(H_n(S_{t,\theta}) = \delta \mid X = (x, t)) \cdot P(X = (x, t))}{P(H_n(S_{t,\theta}) = \delta)}$$

Which for compactness we will write as:

$$P_{X|H_n}((x, t), \delta) = \frac{P_{H_n|X}(\delta, (x, t)) \cdot P_X((x, t))}{P_{H_n}(\delta)} \quad (9)$$

We can easily determine one of these terms:

$$P_X((x, t)) = \prod_{0 \leq n < q} \phi_n(x, t) \quad (10)$$

Because the entropy 5 is calculated as a sum of the entropy of each observation, we can determine the conditional probability  $P(H_n(S_{t,\theta}) = \delta | X = (x, t))$  by subtracting the entropy of  $H_n((x, t))$  from the value of  $\delta$ , and then using our correlation estimate  $\varphi_n$  to predict the probability of  $\delta - H_n((x, t))$ . We include the expanded form of  $P(H_n(S_{t,\theta}))$  to clarify the cancellation of terms:

$$\begin{aligned} P_{H_n}(\delta) &= \frac{|y : H_n(y) = \delta|}{|y : H_n(y) \neq \delta|} \\ P_{H_n|X}(\delta, S_{t,\theta}) &= \frac{|y : H_n(y) = \delta + H_n((x, t))|}{|y : H_n(y) \neq \delta|} \\ \frac{P_{H_n|X}(\delta, S_{t,\theta})}{P_{H_n}(\delta)} &= \frac{|y : H_n(y) = \delta + H_n((x, t))| \cdot |y : H_n(y) \neq \delta|}{|y : H_n(y) = \delta| \cdot |y : H_n(y) \neq \delta|} \\ &= \frac{P_{H_n}(\delta + H_n((x, t)))}{P_{H_n}(\delta)} \end{aligned}$$

Where  $|\cdot|$  denotes the number of elements in the set  $(\cdot)$ . The derivation algorithm can easily be extended to multiple windows by using joint entropy probabilities. Because we do not know anything about the conditional relationships between estimates, we must assume they are i.i.d., and that their joint probability is equal to the sum of the marginal probabilities.

$$P(A|B \cap C) = \frac{P(B \cap C|A) \cdot P(A)}{P(B \cap C)}$$

$$\begin{aligned} P_{X|H_n \cap H_m}((x, t), \delta_u, \delta_v) &= \frac{P_{H_n \cap H_m|X}(\delta_u, \delta_v, (x, t)) \cdot P_X((x, t))}{P_{H_n \cap H_m}(\delta_u, \delta_v)} \\ P_{H_n \cap H_m|X}(\delta_u, \delta_v, (x, t)) &= \prod_{\substack{i=[n,m] \\ j=[u,v]}} P_{H_i}(\delta_j + H_i((x, t))) \\ &= \prod_{\substack{i=[n,m] \\ j=[i,j]}} \varphi_i(\delta_j + H_i((x, t))) \end{aligned} \quad (11)$$

$$\begin{aligned} P_{H_n \cap H_m}(\delta_u, \delta_v) &= \prod_{\substack{i=[n,m] \\ j=[u,v]}} P_{H_i}(\delta_j) \\ &= \prod_{\substack{i=[n,m] \\ j=[u,v]}} \varphi_i(\delta_j) \end{aligned} \quad (12)$$

Which gives us the following as our prediction algorithm:

$$\begin{aligned}\psi_P((x, t) : S_{t,\theta}, \Phi) &\mapsto \mathbb{P} = \frac{\prod_{0 \leq n < q} \varphi_n(H_n(S_{t,\theta})) \cdot \prod_{0 \leq n < q} \phi_n(x, t)}{\prod_{0 \leq n < q} \varphi_n(H_n(S_{t,\theta}) + H_n((x, t)))} \\ &= \prod_{0 \leq n < q} \phi_n(x, t) \cdot \frac{\varphi_n(H_n(S_{t,\theta}))}{\varphi_n(H_n(S_{t,\theta}) + H_n((x, t)))} \quad (13)\end{aligned}$$

## 2.5 Multiple Channel Estimation

Algorithms for determining probability density functions tend to scale exponentially with the dimensionality of the input data. For this reason it would be helpful if the algorithm could operate on independent channels of data and only calculate relationships between channels in cases where the channel's probability distribution is conditional on such relationships.

Extending the existing theory to this situation, we return to our original setting of the problem:

$$\begin{aligned}X^n : \Omega &\mapsto \Omega^n \in \mathbb{R}^d \times \mathbb{R}_+ \\ X^n &= (\Omega^n, \mathcal{F}^n, \mathcal{P}^n) \\ \mathbf{X} &= [X^0, \dots, X^c] \\ C^A &= [\Omega^A, \mathcal{F}^A, \mathcal{P}^A]\end{aligned}$$

We will now refer to observations and estimates using superscript to denote their channel. We previously developed three generalized algorithms,  $\psi_E, \psi_C, \psi_D$ ; we will now extend each one to handle multiple channels.

The estimation algorithm is unchanged in the context of multiple channels. Each channel generates independent estimates over independent time windows. The estimation process is intended to give each channel an understanding of its own behavior at a specific time window, and as such the generation of estimates does not rely on previous behaviors of the channel or the behavior of other channels. We therefore re-write the estimation algorithm to reflect the new notation:

$$S_{t,\theta}^A = \left[ \left( x_i^A, \frac{t_i - t}{\theta} \right) \mid x_i^A \in X^A, t \leq t_i < t + \theta \right] \quad (14)$$

$$\psi_E(S_{t,\theta}^A : \alpha_E^A) \mapsto \phi_n^A \simeq \mathcal{P}^A(t, \theta) \quad (15)$$

$$\Phi^A = [\phi_0^A, \dots, \phi_p^A]$$

$$\Phi = [\Phi^0, \dots, \Phi^c]$$

## 2.6 Multiple Channel Correlation

The correlation algorithm is where the most substantial changes must be made to accomodate multiple channels. The biggest difference from the single-channel approach is that each channel takes place in an independent abstract space. Recall that a critical component of the prediction algorithm is the evaluation of the entropy 5 of a set of observations given an estimate  $\phi_n$  15. If the probability density 15 operates over a different abstract space  $\Omega$  than the observations  $S_{t,\theta}$  are taken from, we cannot calculate the entropy. In other words, we cannot evaluate the entropy of an estimate from channel  $C^B$  using observations of channel  $C^A$ , so we must restrict our entropy definition to a single channel:

$$H_n^A(S_{t,\theta}^A) \mapsto \delta_{\phi_n,t,\theta}^A = \sum_{(x,t) \in S_{t,\theta}^A} -\phi_n^A(x,t) \log \phi_n^A(x,t) \quad (16)$$

In order to determine the conditional relations between channels, we begin with the observation that all channels share the time dimension  $t$  as part of their abstract space  $\Omega^A$ . This allows us to specify a uniform time window  $(t, \theta)$  for all channels  $C^A \in \mathbf{C}$ , and to then evaluate the entropy 16 of each estimate  $\phi_n^A \in \Phi^A$  given the subset of each channel's observations  $S_{t,\theta}^A$  for the time window. Given the time window  $(t, \theta)$ , we can treat these entropy measurements as a coherent set and interpret them as a multi-dimensional random vector in much the same way we treated different time windows in the single-channel case 6.

$$\Delta_{t,\theta} = \{A_{\phi_n} : \delta_{\phi_n,t,\theta}^A \mid \phi_n \in \Phi, 0 \leq A < c\}$$

Where  $\Delta_{t,\theta}\{A_{\phi_n}\} = \delta_{\phi_n,t,\theta}^A$ . Using the same algorithm  $\psi_C$  as we used previously, we can estimate the probability density  $\varphi$  of this random vector.

$$\begin{aligned} \Delta &= [\Delta_{t,\theta} \mid (t, \theta) \in \mathcal{T}] \\ \psi_C(\Delta, \alpha_C) &\mapsto \varphi \\ \varphi(\Delta_{t,\theta}) &\simeq \bigcap_{\substack{\phi_n \in \Phi \\ 0 \leq A < c}} P(H_n^A = \Delta_{t,\theta}\{A_{\phi_n}\}) \end{aligned} \quad (17)$$



## 2.7 Multiple Channel Prediction

We begin by considering the situation in which we wish to make predictions on channel  $C^A$  using information from another,  $C^B$ . The derivation algorithm therefore produces predictions of  $X^A$  over a time window  $(t, \theta)$  based on a set of observations which occur in that time window  $S_{t,\theta}^B$ . To do so we select estimates in both channels:

$$\begin{aligned} C^A &\rightarrow [\phi_0^A, \dots, \phi_p^A] \\ C^B &\rightarrow [\phi_0^B, \dots, \phi_q^B] \end{aligned}$$

We calculate the entropy of each estimate in its own context:

$$\begin{aligned} [\delta_{\phi_0}^A, \dots, \delta_{\phi_p}^A] &= [H_0^A(S_{t,\theta}^A), \dots, H_p^A(S_{t,\theta}^A)] \\ [\delta_{\phi_0}^B, \dots, \delta_{\phi_q}^B] &= [H_0^B(S_{t,\theta}^B), \dots, H_q^B(S_{t,\theta}^B)] \end{aligned}$$

Given an two estimates from the channels, we calculate the probability of each point in  $X^A$  based on the joint probability of the entropy values  $P(\delta^A \cap \delta^B)$ .

$$P_{X|H_n \cap H_m}^A((x, t), \delta^A, \delta^B) = \frac{P_{H_n \cap H_m|X}^A(\delta^A, \delta^B, (x, t)) \cdot P_X^A((x, t))}{P_{H_n \cap H_m}^A(\delta^A, \delta^B)} \quad (18)$$

This time, we can easily determine two of these terms:

$$P_X^A((x, t)) = \prod_{0 \leq i < p} \phi_i^A(x, t) \quad (19)$$

$$P_{H_n \cap H_m}^A(\delta^A, \delta^B) = \varphi(\{A_{\phi_n} : \delta^A, B_{\phi_m} : \delta^B\}) \quad (20)$$

Which leaves us with the process of calculating the effect of  $X^A = (x, t)$  on the entropy of  $\phi_n^A$  and  $\phi_m^B$ . We adopt the same process used in the single-channel case, replacing  $\delta^A$  with  $\delta^A + H_n^A((x, t))$ , and recalculating the joint probability. Since  $\phi_m^B$  is in a different abstract space, we cannot make any useful statements about the dependence of its entropy given a point outside the abstract space being predicted. We can now determine the third of the three terms:

$$\begin{aligned} P_{H_n \cap H_m|X}^A(\delta^A, \delta^B, (x, t)) &= P_{H_n \cap H_m}^A(\delta^A + H_n^A((x, t)), \delta^B) \\ &= \varphi(\{A_{\phi_n} : \delta^A + H_n^A((x, t)), B_{\phi_m} : \delta^B\}) \end{aligned} \quad (21)$$

This result can easily be extended to an arbitrary number of channels and estimates:

$$P_{\Delta}^A(\Delta_{t,\theta}) = \varphi(\Delta_{t,\theta}) \quad (22)$$

$$P_{\Delta|X}^A(\Delta_{t,\theta}, (x, t)) = \varphi(\{A_{\phi_n} : \delta_{\phi_n}^A + H_n^A((x, t)), \Delta_{t,\theta} \wedge A_{\phi_n}\}) \quad (23)$$

Which gives us the final prediction equation:

$$\begin{aligned} \psi(x, t) &\mapsto \mathbb{P} = P(X^A = (x, t)) \\ &= \prod_{0 \leq n < p} \phi_n^A(x, t) \cdot \frac{\varphi(\{A_{\phi_n} : \delta_{\phi_n}^A + H_n^A((x, t)), \Delta_{t,\theta} \wedge A_{\phi_n}\})}{\varphi(\Delta_{t,\theta})} \end{aligned} \quad (24)$$

### 3 Sample Implementation

Implementing the inference architecture requires the specification of the two density estimation functions,  $\psi_E$  and  $\psi_C$ . We will use the same algorithm for both; Vapnik's SVM for conditional density estimation. The SVM algorithm has been chosen due to its simplicity of parameters and determinacy.

#### 3.1 System Architecture

The first challenge is the selection of  $\mathcal{T}$ , the intervals used to create the estimates  $\Phi$ . I think the best option here is to treat each observation as a  $t_n$ , and to determine  $\theta$  as part of the selection process, in which case each observation defines a subset  $S_{t,\theta}$  for each value of  $\theta$ . The challenge is therefore reduced to selecting a subset of observations to use as the estimate set.

I think a simple method of doing this is to define an SVM over the inputs which minimizes the entropy of a neighborhood of observations relative to a PDF generated by the SVM, with constraints designed to minimize the number of neighborhoods used in the mapping. In this case the mapping would be accomplished through the weighted sum of a set of neighborhoods, and the challenge is to select a sparse yet expressive set of neighborhoods. My suspicion is that this problem is virtually identical to the PDF problem - in both cases we're minimizing the square loss, except in this case instead of using the CDF for  $Y$ , we a subset of  $X$  (at each observation we use a subset of the observations offset so the subset is centered at  $X.t = 0$ )

The trick to making this all work is the kernel function. The kernel function needs to implement both  $\psi_E$  and  $H_n(S)$ . It needs to take two sets of observations and determine their correlation. In the same way we used multiple kernels

with different width parameters, we might be able to use multiple kernels with different  $\theta$  values in this case.

This method may be able to solve the problem without explicitly generating any density estimates, in which case we \*might\* be able to do something simple like Parzen windows on the SVM output (since in this case sparsity isn't so important; sparsity is accomplished through the SVM). This would require finding a substitute for  $\psi_E$  and  $H_n(S)$  capable of implementing the same behavior. Specifically, the kernel would need to evaluate the similarity between two sets of observations.

### 3.2 SVM-Density Algorithm

Our objective is to determine the conditional density function and regression functions

$$p(y|x) = \sum_{i=1}^{\ell} \beta_i K_{\gamma}(x, x_i) K_{\gamma}(y, y_i), \beta \geq 0 \quad (25)$$

$$r(x) = \sum_{i=1}^{\ell} y_i \beta_i K_{\gamma}(x, x_i) \quad (26)$$

Which are determined by optimizing for the values of  $\beta_i$  which minimize

$$W_{\gamma}(\beta) = \Omega(f) = \sum_{i,j=1}^{\ell} \beta_i \beta_j K_{\gamma}(x_j, x_i) K_{\gamma}(y_j, y_i) \quad (27)$$

Subject to the Constraints

$$\max_p \left| \sum_{i=1}^{\ell} \beta_i \frac{1}{\ell} \sum_{j=1}^{\ell} K_{\gamma}(x_j, x_i) \theta(x_p - x_j) \int_{-\infty}^{y_p} K_{\gamma}(y', y_i) dy' - F_{\ell}(x_p, y_p) \right| = \sigma_{\ell} \quad (28)$$

$$\sum_{i=1}^{\ell} \left( \frac{1}{\ell} \sum_{j=1}^{\ell} \beta_j K_{\gamma}(x_i, x_j) \right) = 1 \quad (29)$$

We use a kernel function defined as follows, where  $\gamma$  is chosen from the admissible set to control the desired properties of the solution.

$$K_{\gamma}(x, x_i) = a(\gamma) \exp \left( - \left| \frac{x - x_i}{\gamma} \right|^p \right), 0 < p \leq 2 \quad (30)$$

We choose  $\sigma$  based on an appropriate quantile of the universal distribution of discrepancy:

$$\sigma = \frac{q}{\sqrt{\ell}} \quad (31)$$

### **3.3 Data Pre-Processing**

### **3.4 Parameter Selection**

## **4 Results**

The architecture has been tested against several data sets. In all cases the system parameters are left unchanged to eliminate the possibility of optimizing the system performance to best match the known outcomes.

### **4.1 Eunite Competition Data**

### **4.2 Santa Fe Data**

### **4.3 CATS Benchmark Data**

### **4.4 Results Summary**

## **5 Further Research**

The system as described thus far makes several assumptions for simplicity which would lead to unnecessary computational demands. We now spend some time discussing methods of reducing the computational demands of the system.

### **5.1 Data Retention**

### **5.2 Estimation Optimization**

The generation of estimates has been described as a process of selecting time windows at random from the full set of training data. Selecting windows at random is not necessary to satisfy the i.i.d. requirements of producing an accurate probability estimate, so long as the selection of time windows for which entropy is calculated is random. Let's consider the characteristics of a 'good' set of estimates for prediction.

The most obvious characteristic of a 'good' estimate is that it has a low average entropy, which is to say that it frequently captures the behavior of observed data. If an estimate is not applicable to the observed data, its influence on predictions will be marginal, and the computational time required to calculate its entropy and influence on a prediction will have been wasted.

Another characteristic of a 'good' set of estimates is that they have minimal redundancy. Again, if two estimates are essentially identical, their influence on prediction tasks will be identical, and the computational demands of evaluating their entropy and influence will be wasted.

Unfortunately, these two parameters will likely be mutually exclusive, so some method of balancing them is required.

### 5.3 Correlation

It is critical that correlation takes place using i.i.d. data in order to produce accurate probability estimates. We had previously assumed that each time window used to generate estimates would be used to correlate the estimates, however it is not necessary that the time windows used to determine the entropy values used to correlate estimates correspond to the time windows used to generate estimates. This allows us to select a number of time windows to use which balances the computational demands of density estimation with the need for accuracy.

### 5.4 Prediction

For prediction, it is not necessary to actually use each of the estimates in generating predictions, so long as we have accurate estimates of their individual probabilities and the joint probability of any two estimates. Since we must evaluate a number of joint probabilities equal to the square of the number of estimates being considered, selecting a useful subset of the defined estimates can provide substantial reductions in computational demands.

### 5.5 Ensemble System

### 5.6 Prediction with Heterogenous Time Windows

## 6 Conclusion

Eat it, bitches